# VitisOmics

(see contributors below)

October 28, 2015

# Contents

# 1 Overview

This document describes the "VitisOmics" project. This project aims at handling "omics" data in the genus Vitis (e.g. grapevine) in an open and reproducible way. Nevertheless, it requires some basic knowledge and skills about bioinformatics on GNU/Linux computers.

A large amount of such "omics" data are already available, and several committes from the IGGP (International Grape Genome Program) strive at improving interoperability. However, several issues remain, among which:

- partially-overlapping data present at multiple locations (URGI, CRIBI, NCBI, EBI, etc);

- data downloadable as files in various formats not always easy to interchange (fasta, genbank, gff3, gtf, etc);

- data often available without meta-data inside the file;

- large files not always available in compressed form;

- main efforts dedicated to wet-labs grapevine biologists.

These have consequences in terms of ambiguity, inefficiency, potential mistakes, etc. Therefore, I hope that my attempt, via the usage of git and GitHub, could prove for the community to be a useful addition to the IGGP efforts.

The repository can be easily cloned from GitHub:

```
git clone https://github.com/timflutre/VitisOmics.git
```

The project directory is organized as advised by Noble (PLoS Computational Biology 2009). On any Unix-like system, it can be easily compressed and transferred (ignoring large data files):

```
cd ..; tar -czvf VitisOmics.tar.gz \
--exclude=VitisOmics/data --exclude=VitisOmics/results \
```

```
--exclude="*~" --exclude=".*" VitisOmics
```

In order to concretely promote collaborative editing in a distributed manner, the content of the "VitisOmics" repository should be based on plain text files. As a consequence, this document is written in the org format, and can thus be automatically exported, best by emacs, but also by pandoc, into the pdf and html formats for easy reading. A choice is also made to use as much as possible softwares widely available on any GNU/Linux computer, such as bash, awk, etc, but of course it may sometimes be much easier to use R (with Bioconductor), Python (with Biopython), etc.

Last but not least, feel free to contribute, by reporting issues or forking the repository!

## 1.1 Contributors

The person roles comply with R's guidelines (The R Journal Vol. 4/1, June 2012).

- Timothée Flutre (cre,aut)
- Gautier Sarah (ctb)
- …

## 1.2 References

The NCBI has web pages about genome assembly:

- Assembly information;
- NCBI Genome Assembly Model;
- AGP specification.

The original article for grapevine (PN40024) is a must read:

- Jaillon, et al. The grapevine genome sequence suggests ancestral hexaploidization in major angiosperm phyla. Nature 449, 463-467 (2007).

About genome annotations:

- Grimplet et al. Comparative analysis of grapevine whole-genome gene predictions, functional annotation, categorization and integration of the predicted gene sequences. BMC Research Notes 5, 213+ (2012).
- Grimplet et al. The grapevine gene nomenclature system. BMC Genomics 15, 1077+ (2014).

- Vitulo et al. A deep survey of alternative splicing in grape reveals changes in the splicing machinery related to tissue, stress condition and genotype. BMC Plant Biology 14, 99+ (2014).

# 2 Data

```
mkdir -p data; cd data/
```

TODO: retrieve genome data from other cultivars than PN40024, e.g. Sultanina

## 2.1 URGI

- `https://urgi.versailles.inra.fr/Species/Vitis`

```
mkdir -p urgi; cd urgi/
../../src/download_urgi.bash
```

Remarks:

- when needed, the script decompresses zip files and compress them again but with gzip instead;
- the AGP file for the PN40024 8x assembly is not available at URGI's web site.

At the bottom of the page, one can find the following assemblies summary.

12x.2 assembly:

- contigs: 14642
- scaffolds (> 2kb): 2059
- mapped scaffolds: 366 (coverage of the genome: 95%)

12X.0 assembly:

- contigs: 14642
- scaffolds (> 2kb): 2059
- mapped scaffolds: 211 (coverage of the genome: 91.2%)

8X assembly:

- contigs: 19577
- scaffolds: 3514
- mapped ultracontigs: 191 (coverage of the genome: 68.9%)

N. Choisne from URGI (personal communication, 13/10/2015):

- "mapped scaffolds": scaffolds anchored on the linkage groups (i.e. chromosomes) using the markers from the reference genetic map;

- unmapped scaffolds hence are unanchored, and gathered into chrUn;

- "supercontig" is a synonym of "scaffold";

- "ultracontig": one level above supercontigs; localized and orientated according to data from BAC libraries.

## 2.2 NCBI

- http://www.ncbi.nlm.nih.gov/genome/401

- ftp://ftp.ncbi.nlm.nih.gov/genomes/Vitis_vinifera/

```
mkdir -p ncbi; cd ncbi/
../../src/download_ncbi.bash
```

Remarks:

- the important file scaffold_names provides the correspondence between original scaffold names (i.e. from the sequencing center) and various NCBI identifiers (RefSeq, GenBank, etc);

- in ARCHIVE/, BUILD.1.1/ corresponds to the 8x genome sequences of PN40024.

## 2.3 EBI

```
mkdir -p ebi; cd ebi/
../../src/download_ebi.bash
```

Remarks:

- a genome soft-masked by RepeatMasker is available.

## 2.4 CRIBI

- http://genomes.cribi.unipd.it/grape/

```
mkdir -p cribi; cd cribi/
../../src/download_cribi.bash
```

## 2.5 Genoscope

- http://www.genoscope.cns.fr/externe/Download/Projets/Projet_ML/data/

```
mkdir -p genoscope; cd genoscope/
../../src/download_genoscope.bash
```

# 3 Results

```
mkdir -p results; cd results/
```

TODO: compress fasta files with `bgzip` instead of `gzip`

## 3.1 Comparisons of original "assembly" files

Files from URGI:

```
cd urgi/
zcat VV_8X_embl_98_WGS_contigs.fsa.gz | grep -c ">" # 19577
zcat VV_8X_embl_98_Scaffolds.fsa.gz | grep -c ">" # 3514
zcat VV_chr8x.fsa.gz | grep -c ">" # 35
zcat VV_12X_embl_102_WGS_contigs.fsa.gz | grep -c ">" # 14642
zcat VV_12X_embl_102_Scaffolds.fsa.gz | grep -c ">" # 2059
zcat VV_chr12x.fsa.gz | grep -c ">" # 33
cat 12x0_chr.agp | wc -l # 390
cat 12x0_scaffolds.lg | wc -l # 2059
cat 12x0_chr.lg | wc -l # 33
zcat 12Xv2_grapevine_genome_assembly.fa.gz | grep -c ">" # 20
```

Files from NCBI:

```
cd ncbi/
ls ARCHIVE/BUILD.1.1/CHRS/vvi_ref_chr*.fa.gz | grep -v "Pltd" | while read f; do
    zcat $f; done | grep -c ">" # 3514
ls ARCHIVE/BUILD.1.1/Assembled_chromosomes/vvi_ref_chr*.fa.gz | while read f; do
    zcat $f; done | grep -c ">" # 19
zcat ARCHIVE/BUILD.1.1/allcontig.agp.gz | grep -v "#" | cut -f 5 | sort | uniq -c #
    F=1 N=16063 W=19577
cat ARCHIVE/BUILD.1.1/scaffold_names | sed 1d | wc -l # 3514
ls CHRS/vvi_ref_12X_chr*.fa.gz | grep -v -E "Pltd|MT" | while read f; do zcat $f;
    done | grep -c ">" # 2059
```

```
ls Assembled_chromosomes/vvi_ref_12X_chr*.fa.gz | grep -v -E "Pltd|MT" | while read
    f; do zcat $f; done | grep -c ">" # 19
cat scaffold_names | sed 1d | wc -l # 2061
```

See also the script `src/vitisomics.R` using R and Bioconductor. It confirms that the 12x scaffolds have the exact same sequence, whether they come from the URGI or the NCBI. Note however that the file from the NCBI allows to know easily on which chromosome a placed sequences is.

Remarks concerning PN40024 at URGI:

- the file `12x0_chr.agp.info` doesn't correspond to `12x0_chr.agp` (it doesn't even correspond to the description of a proper AGP file, as specified here);

- no mitochondrial nor chloroplastic data are available.

Remarks concerning PN40024 at NCBI:

- contig `NC_007957.1` in `ARCHIVE/BUILD.1.1/allcontig.agp.gz` (with `fragment_type=F` for "finished") corresponds to the chloroplast;

- `scaffold_names` contains all 2059 scaffolds of nuclear DNA as well as the assembled genome of the mitochondria and the chloroplast.

For its build 1.1 (corresponding to the 8x sequences of the PN40024 variety), the NCBI has one file per assembled chromosome. However, all unlocalized and unplaced scaffolds are gathered in a single file `chrUn`. This is not the case at URGI which has unlocalized scaffolds in files as `chr3_random` and a `chrUn_random` file with all unplaced scaffolds (and only them). Unfortunately, the NCBI has the annotation of the 8x (in the GenBank format), but the URGI hasn't.

## 3.2 Manipulations of files from URGI

```
mkdir -p urgi; cd urgi/
```

### 3.2.1 Reformat sequence headers for VITVI_PN40024_8x_chroms_URGI

Launch script:

```
ln -s ../../data/urgi/VV_chr8x.fsa.gz .
echo "../../src/reformat_VV_chr8x.bash" \
  | qsub -cwd -j y -V -N reformat_VV_chr8x -q normal.q
```

Check:

```
zcat VV_chr8x.fsa.gz | wc -l # 8291865
zcat VV_chr8x.fsa.gz | grep -c ">" # 35
zcat VITVI_PN40024_8x_chroms_URGI.fa.gz | wc -l # 8291865
zcat VITVI_PN40024_8x_chroms_URGI.fa.gz | grep -c ">" # 35
diff <(zcat VV_chr8x.fsa.gz) <(zcat VITVI_PN40024_8x_chroms_URGI.fa.gz)
```

Only the headers differ, not the sequences, so everything is fine.

Basic stats:

```
zcat VITVI_PN40024_8x_chroms_URGI.fa.gz | md5sum # 4b6ea1cb4ff189ac587fa269077885b5
```

Length of each sequence:

```
zcat VITVI_PN40024_8x_chroms_URGI.fa.gz \
  | awk 'BEGIN{RS=">"} {split($0,a,"\n");
if(length(a)==0)next;
sum=0; for(i=2;i<=length(a);++i){sum+=length(a[i])};
print a[1]": "sum; sumTot+=sum} END{print sumTot}'
```

| header | length (bp) |
| --- | --- |
| chr1 CU462738\|Vitis vinifera\|PN40024\|assembly8x\|chromosome$_1$ | 15630816 |
| chr10 CU462747\|Vitis vinifera\|PN40024\|assembly8x\|chromosome$_{10}$ | 9647040 |
| chr10$_{random}$ | 2206354 |
| chr11 CU462748\|Vitis vinifera\|PN40024\|assembly8x\|chromosome$_{11}$ | 13936303 |
| chr11$_{random}$ | 1958407 |
| chr12 CU462749\|Vitis vinifera\|PN40024\|assembly8x\|chromosome$_{12}$ | 18540817 |
| chr12$_{random}$ | 2826407 |
| chr13 CU462750\|Vitis vinifera\|PN40024\|assembly8x\|chromosome$_{13}$ | 15191948 |
| chr13$_{random}$ | 1580403 |
| chr14 CU462751\|Vitis vinifera\|PN40024\|assembly8x\|chromosome$_{14}$ | 19480434 |
| chr14$_{random}$ | 5432426 |
| chr15 CU462752\|Vitis vinifera\|PN40024\|assembly8x\|chromosome$_{15}$ | 7693613 |
| chr15$_{random}$ | 4297576 |
| chr16 CU462753\|Vitis vinifera\|PN40024\|assembly8x\|chromosome$_{16}$ | 8158851 |
| chr16$_{random}$ | 4524411 |
| chr17 CU462754\|Vitis vinifera\|PN40024\|assembly8x\|chromosome$_{17}$ | 13059092 |
| chr17$_{random}$ | 1763011 |
| chr18 CU462755\|Vitis vinifera\|PN40024\|assembly8x\|chromosome$_{18}$ | 19691255 |
| chr18$_{random}$ | 5949186 |
| chr19 CU462756\|Vitis vinifera\|PN40024\|assembly8x\|chromosome$_{19}$ | 14071813 |
| chr19$_{random}$ | 1912523 |
| chr1$_{random}$ | 5496190 |
| chr2 CU462739\|Vitis vinifera\|PN40024\|assembly8x\|chromosome$_2$ | 17603400 |
| chr2$_{random}$ | 60809 |
| chr3 CU462740\|Vitis vinifera\|PN40024\|assembly8x\|chromosome$_3$ | 10186927 |
| chr3$_{random}$ | 1343266 |
| chr4 CU462741\|Vitis vinifera\|PN40024\|assembly8x\|chromosome$_4$ | 19293076 |
| chr5 CU462742\|Vitis vinifera\|PN40024\|assembly8x\|chromosome$_5$ | 23428299 |
| chr6 CU462743\|Vitis vinifera\|PN40024\|assembly8x\|chromosome$_6$ | 24148918 |
| chr7 CU462744\|Vitis vinifera\|PN40024\|assembly8x\|chromosome$_7$ | 15233747 |
| chr7$_{random}$ | 176143 |
| chr8 CU462745\|Vitis vinifera\|PN40024\|assembly8x\|chromosome$_8$ | 21557227 |
| chr8$_{random}$ | 12125 |
| chr9 CU462746\|Vitis vinifera\|PN40024\|assembly8x\|chromosome$_9$ | 16532244 |
| chrUn$_{random}$ | 154883714 |
| total | 497508771 |

### 3.2.2 Reformat sequence headers for `VITVI_PN40024_12x_v0_scaffolds_EMBL_r102`

Launch script:

```
ln -s ../../data/urgi/VV_12X_embl_102_Scaffolds.fsa.gz .
echo "../../src/reformat_VV_12X_embl_102_Scaffolds.bash" \
  | qsub -cwd -j y -V -N reformat_VV_12X_embl_102_Scaffolds -q normal.q
```

Check:

```
zcat VV_12X_embl_102_Scaffolds.fsa.gz | wc -l # 8091565
zcat VV_12X_embl_102_Scaffolds.fsa.gz | grep -c ">" # 2059
zcat VITVI_PN40024_12x_v0_scaffolds_EMBL_r102.fa.gz | wc -l # 8091565
zcat VITVI_PN40024_12x_v0_scaffolds_EMBL_r102.fa.gz | grep -c ">" # 2059
diff <(zcat VV_12X_embl_102_Scaffolds.fsa.gz) <(zcat
    VITVI_PN40024_12x_v0_scaffolds_EMBL_r102.fa.gz)
```

Only the headers differ, not the sequences, so everything is fine.

Basic stats:

```
zcat VITVI_PN40024_12x_v0_scaffolds_EMBL_r102.fa.gz | md5sum # 4
    fa2432d7a66c019c7cb41ee4d0cb7bc
zcat VITVI_PN40024_12x_v0_scaffolds_EMBL_r102.fa.gz | grep -v ">" | md5sum #
    df5cdb0c6f73cb133261905374cdf2f2
```

### 3.2.3 Reformat sequence headers for `VITVI_PN40024_12x_v0_chroms_URGI`

Launch script:

```
ln -s ../../data/urgi/VV_chr12x.fsa.gz .
echo "../../src/reformat_VV_chr12x.bash" \
  | qsub -cwd -j y -V -N reformat_VV_chr12x -q normal.q
```

Check:

```
zcat VV_chr12x.fsa.gz | wc -l # 8240706
zcat VV_chr12x.fsa.gz | grep -c ">" # 33
zcat VITVI_PN40024_12x_v0_chroms_URGI.fa.gz | wc -l # 8240706
zcat VITVI_PN40024_12x_v0_chroms_URGI.fa.gz | grep -c ">" # 33
diff <(zcat VV_chr12x.fsa.gz) <(zcat VITVI_PN40024_12x_v0_chroms_URGI.fa.gz)
```

Only the headers differ, not the sequences, so everything is fine.

Basic stats:

```
zcat VITVI_PN40024_12x_v0_chroms_URGI.fa.gz | md5sum #
    eff315994fafe35333462b9595e10ce5
```

### 3.2.4 Reformat sequence headers for `VITVI_PN40024_12x_v2_chroms_URGI`

Launch script:

```
ln -s ../../data/urgi/12Xv2_grapevine_genome_assembly.fa.gz .
echo "../../src/reformat_12Xv2_grapevine_genome_assembly.bash" \
  | qsub -cwd -j y -V -N reformat_12Xv2_grapevine_genome_assembly -q normal.q
```

Check:

```
zcat 12Xv2_grapevine_genome_assembly.fa.gz | wc -l # 8103449
zcat 12Xv2_grapevine_genome_assembly.fa.gz | grep -c ">" # 20
zcat VITVI_PN40024_12x_v2_chroms_URGI.fa.gz | wc -l # 8103449
zcat VITVI_PN40024_12x_v2_chroms_URGI.fa.gz | grep -c ">" # 20
diff <(zcat 12Xv2_grapevine_genome_assembly.fa.gz) <(zcat
    VITVI_PN40024_12x_v2_chroms_URGI.fa.gz)
```

Only the headers differ, not the sequences, so everything is fine.

Basic stats:

```
zcat VITVI_PN40024_12x_v2_chroms_URGI.fa.gz | md5sum # 4
    e487c28eaf19ef59b0b6128b73935af
```

Length of each sequence:

```
zcat VITVI_PN40024_12x_v2_chroms_URGI.fa.gz \
  | awk 'BEGIN{RS=">"} {split($0,a,"\n");
if(length(a)==0)next;
sum=0; for(i=2;i<=length(a);++i){sum+=length(a[i])};
print a[1]": "sum; sumTot+=sum} END{print sumTot}'
```

| header | length (bp) |
| --- | --- |
| chr1 Vitis vinifera\|PN40024\|assembly12x.2 | 24233538 |
| chr2 Vitis vinifera\|PN40024\|assembly12x.2 | 18891843 |
| chr3 Vitis vinifera\|PN40024\|assembly12x.2 | 20695524 |
| chr4 Vitis vinifera\|PN40024\|assembly12x.2 | 24711646 |
| chr5 Vitis vinifera\|PN40024\|assembly12x.2 | 25650743 |
| chr6 Vitis vinifera\|PN40024\|assembly12x.2 | 22645733 |
| chr7 Vitis vinifera\|PN40024\|assembly12x.2 | 27355740 |
| chr8 Vitis vinifera\|PN40024\|assembly12x.2 | 22550362 |
| chr9 Vitis vinifera\|PN40024\|assembly12x.2 | 23006712 |
| chr10 Vitis vinifera\|PN40024\|assembly12x.2 | 23503040 |
| chr11 Vitis vinifera\|PN40024\|assembly12x.2 | 20118820 |
| chr12 Vitis vinifera\|PN40024\|assembly12x.2 | 24269032 |
| chr13 Vitis vinifera\|PN40024\|assembly12x.2 | 29075116 |
| chr14 Vitis vinifera\|PN40024\|assembly12x.2 | 30274277 |
| chr15 Vitis vinifera\|PN40024\|assembly12x.2 | 20304914 |
| chr16 Vitis vinifera\|PN40024\|assembly12x.2 | 23572818 |
| chr17 Vitis vinifera\|PN40024\|assembly12x.2 | 18691847 |
| chr18 Vitis vinifera\|PN40024\|assembly12x.2 | 34568450 |
| chr19 Vitis vinifera\|PN40024\|assembly12x.2 | 24695667 |
| chrUkn Vitis vinifera\|PN40024\|assembly12x.2 | 27389308 |
| total | 486205130 |

### 3.2.5  Format `VITVI_PN40024_12x_v0_chroms_URGI` for **BLAST**

TODO: change Vvin to VITVI Launch:

```
echo "../../src/blast_format.bash VITVI_PN40024_12x_v0_chroms_URGI.fa.gz" \
  | qsub -cwd -j y -V -N blast_format_VITVI_PN40024_12x_v0_chroms_URGI -q normal.q
```

### 3.2.6  Index `VITVI_PN40024_12x_v0_chroms_URGI` for **BWA**

Launch:

```
echo "../../src/bwa_index_VITVI_PN40024_12x_v0_chroms_URGI.bash" \
  | qsub -cwd -j y -V -N bwa_index_VITVI_PN40024_12x_v0_chroms_URGI -q normal.q
```

### 3.2.7 Index VITVI_PN40024_12x_v2_chroms_URGI for BWA

Launch:

```
echo "../../src/bwa_index_VITVI_PN40024_12x_v2_chroms_URGI.bash" \
  | qsub -cwd -j y -V -N bwa_index_VITVI_PN40024_12x_v2_chroms_URGI -q normal.q
```

### 3.2.8 Prepare VITVI_PN40024_12x_v2_chroms_URGI for SAMtools and Picard

Make an index as well as a SAM header.

Launch:

```
echo "../../src/samtools-picard_prep_VITVI_PN40024_12x_v2_chroms_URGI.bash" \
  | qsub -cwd -j y -V -N samtools-picard_prep_VITVI_PN40024_12x_v2_chroms_URGI -q
     normal.q
```

### 3.2.9 Index VITVI_PN40024_12x_v0_chroms_URGI for Bowtie2

Launch:

```
echo "../../src/bowtie2_index_VITVI_PN40024_12x_v0_chroms_URGI.bash" \
  | qsub -cwd -j y -V -N bowtie2_build_VITVI_PN40024_12x_v0_chroms_URGI -q normal.q
```

### 3.2.10 Index VITVI_PN40024_12x_v2_chroms_URGI for Bowtie2

Launch:

```
echo "../../src/bowtie2_index_VITVI_PN40024_12x_v2_chroms_URGI.bash" \
  | qsub -cwd -j y -V -N bowtie2_build_VITVI_PN40024_12x_v2_chroms_URGI -q normal.q
```

### 3.2.11 Index VITVI_PN40024_12x_v2_chroms_URGI for Bowtie2 compatible with Tassel

Tassel requires numbers as chromosome identifiers.

Launch:

```
echo "../../src/bowtie2_index_VITVI_PN40024_12x_v2_chroms_URGI_for_Tassel.bash" \
  | qsub -cwd -j y -V -N bowtie2_build_VITVI_PN40024_12x_v2_chroms_URGI_for_Tassel -
     q normal.q
```

### 3.2.12 Translate CRIBI annotations from 12X.0 to 12X.2

Requirement: use or write a script taking as input the 12X.0 GFF3 file as well as the 12.0-12.2 AGP file, and returns as output the 12X.2 GFF3 file

The URGI provides the following AGP file: `golden_path_V2_111113_allChr.csv`. Unfortunately, after looking at the official specification of the AGP format, the URGI file doesn't seem to be valid, neither for version 1.1, nor 2.2. After contacting URGI, they told me they were working on it (October 2015).

Another script was developped by G. Sarah, but it suffers from several issues.

TODO: test CrossMap

## 3.3 Manipulations of files from NCBI

```
mkdir -p ncbi; cd ncbi/
```

### 3.3.1 Reformat sequence headers for `VITVI_PN40024_8x_scaffolds_NCBI`

Launch script:

```
ls ../../data/ncbi/ARCHIVE/BUILD.1.1/CHRS/vvi_ref_chr*.fa.gz | grep -v -E "Pltd" |
    while read f; do ln -s $f .; done
echo "../../src/reformat_scaffs_NCBI-8x.bash" \
  | qsub -cwd -j y -V -N reformat_scaffs_NCBI-8x -q normal.q
```

Check:

```
\ls vvi_ref_chr* | while read f; do zcat $f; done | wc -l # 6963886
\ls vvi_ref_chr* | while read f; do zcat $f; done | grep -c ">" # 3514
zcat VITVI_PN40024_8x_scaffolds_NCBI.fa.gz | wc -l # 6963886
zcat VITVI_PN40024_8x_scaffolds_NCBI.fa.gz | grep -c ">" # 3514
diff <(\ls -v vvi_ref_chr* | while read f; do zcat $f; done) <(zcat
    VITVI_PN40024_8x_scaffolds_NCBI.fa.gz)
```

Only the headers differ, not the sequences, so everything is fine.

Basic stats:

```
zcat VITVI_PN40024_8x_scaffolds_NCBI.fa.gz | md5sum #
    a66f86ab2d89eb582935454ae3b7a49d
```

### 3.3.2 Reformat sequence headers for `VITVI_PN40024_12x_v0_scaffolds_NCBI`

Launch script:

```
ls ../../data/ncbi/CHRS/vvi_ref_12X_chr*.fa.gz | grep -v -E "Pltd|MT" | while read f
    ; do ln -s $f .; done
echo "../../src/reformat_scaffs_NCBI-12x.bash" \
  | qsub -cwd -j y -V -N reformat_scaffs_NCBI-12x -q normal.q
```

Check:

```
\ls vvi_ref_12X_chr* | while read f; do zcat $f; done | wc -l # 6934292
\ls vvi_ref_12X_chr* | while read f; do zcat $f; done | grep -c ">" # 2059
zcat VITVI_PN40024_12x_v0_scaffolds_NCBI.fa.gz | wc -l # 6934292
zcat VITVI_PN40024_12x_v0_scaffolds_NCBI.fa.gz | grep -c ">" # 2059
diff <(\ls -v vvi_ref_12X_chr* | while read f; do zcat $f; done) <(zcat
    VITVI_PN40024_12x_v0_scaffolds_NCBI.fa.gz) | less
```

Only the headers differ, not the sequences, so everything is fine.

Basic stats:

```
zcat VITVI_PN40024_12x_v0_scaffolds_NCBI.fa.gz | md5sum # 20
    fa822ed5679519a20fe768c422a701
zcat VITVI_PN40024_12x_v0_scaffolds_NCBI.fa.gz | grep -v ">" | md5sum # 9
    ddbb5761fe0e4356c7ef73410011ccb
```

### 3.3.3 Format `VITVI_PN40024_8x_scaffolds_NCBI` for BLAST

Launch:

```
echo "../../src/blast_format.bash VITVI_PN40024_8x_scaffolds_NCBI.fa.gz" \
  | qsub -cwd -j y -V -N blast_format_VITVI_PN40024_8x_scaffolds_NCBI -q normal.q
```

### 3.3.4 Format `VITVI_PN40024_12x_v0_scaffolds_NCBI` for BLAST

Launch:

```
echo "../../src/blast_format.bash VITVI_PN40024_12x_v0_scaffolds_NCBI.fa.gz" \
  | qsub -cwd -j y -V -N blast_format_VITVI_PN40024_12x_v0_scaffolds_NCBI -q normal.
    q
```

### 3.3.5 Convert `gbs` files to `GFF3`

Check that there is one `LOCUS` entry per scaffold:

```
ls ../../data/ncbi/ARCHIVE/BUILD.1.1/CHRS/vvi_ref_chr*.gbs.gz | grep -v "Pltd" |
    while read f; do zcat $f; done | grep -c "LOCUS" # 3514
```

Use the bp$_{\text{genbank2gff3.pl}}$ script from BioPerl:

```
zcat ../../../data/ncbi/ARCHIVE/BUILD.1.1/CHRS/vvi_ref_chr1.gbs.gz | bp_genbank2gff3
    .pl -in stdin -out stdout | gzip > vvi_ref_chr1.gff3.gz
# Error::throw
# Bio::Root::Root::throw /usr/local/share/perl5/Bio/Root/Root.pm:449
# Bio::SeqFeature::Tools::Unflattener::unflatten_seq /usr/local/share/perl5/Bio/
    SeqFeature/Tools/Unflattener.pm:1636
# main::unflatten_seq /usr/local/bin/bp_genbank2gff3.pl:1030
# /usr/local/bin/bp_genbank2gff3.pl:504
```

Use the convert$_{\text{genbanktogff3.py}}$ script from biocode:

```
zcat ../../data/ncbi/ARCHIVE/BUILD.1.1/CHRS/vvi_ref_chr1.gbs.gz > vvi_ref_8x_chr1.
    gbs
convert_genbank_to_gff3.py -i vvi_ref_8x_chr1.gbs -o vvi_ref_8x_chr1.gff3 --no_fasta
# File "convert_genbank_to_gff3.py", line 196, in <module> main()
# File "convert_genbank_to_gff3.py", line 95, in main
# locus_tag = feat.qualifiers['locus_tag'][0]
# KeyError: 'locus_tag'
```

Additional remarks:

- it is written in Python;

- it uses Biopython, but also custom libraries;

- it is on GitHub;

- it doesn't handle gzipped file as input;

- it skips features not from type gene, mRNA, tRNA, rRNA and CDS.

Use the `GFF` library from BCBio (not yet integrated into Biopython) as explained here:

```
zcat ../../data/ncbi/ARCHIVE/BUILD.1.1/CHRS/vvi_ref_chr1.gbs.gz > vvi_ref_8x_chr1.
    gbs
genbank_to_gff.py vvi_ref_8x_chr1.gbs
```

Remarks:

- the `sequence-region` are interspersed in the output file;
- what does the first data line correspond to, with source `annotation`?
- the source is present in the output as a feature;
- why is =feature = added in the 2nd field?
- why is it written `db_xref` instead of `Dbxref` (from official specification)?
- same for `note` instead of `Note`?
- exons seem to have 2nd field as `feature mRNA`

TODO: Use gffutils (doc, code)

TODO: Use a custom script based on Biopython only:

```
genbank2gff3.py -i ../../data/ncbi/ARCHIVE/BUILD.1.1/CHRS/vvi_ref_chr1.gbs.gz -o
    vvi_ref_8x_chr1.gff.gz -t 29760 -g "NCBI 1.1" -s Genbank
```

TODO: check for "pseudo" but empty

## 3.4   Creation of R/Bioconductor packages

- `http://www.bioconductor.org/`
- Huber, W. et al. Orchestrating high-throughput genomic analysis with bioconductor. Nature Methods 12, 115-121 (2015). URL `http://dx.doi.org/10.1038/nmeth.3252`.

TODO: see AnnotationHub

### 3.4.1   BSgenome IGGP12Xv2 package

`http://bioconductor.org/packages/release/bioc/html/BSgenome.html`

Retrieve the sequence data from URGI:

```
cd results/
mkdir -p make_BSgenome_IGGP12Xv2
cd make_BSgenome_IGGP12Xv2/
ln -s ../../data/urgi/12Xv2_grapevine_genome_assembly.fa.gz .
```

Split into one chromosome per file (in the headers, discard everything after the first space):

```
zcat 12Xv2_grapevine_genome_assembly.fa.gz | awk 'BEGIN{RS=">"} {if(NF==0)next;
    split($0,a,"\n"); split(a[1],b," "); print b[1]; print ">"b[1] > b[1]".fa"; for(
    i=2;i<length(a);++i){print a[i] >> b[1]".fa"}}'
gzip chr*.fa
```

Using the latest version of Bioconductor and its BSgenome package, prepare the seed file (`IGGP12Xv2_seed.txt`) by hand as indicated in the vignette as well as in the official R manual "Writing R extensions". Following this article, I chose the CC0 license (present in the R list of licenses in `share/licenses/license.db`). Following suggestions from Hervé Pagès (Bioconductor staff):

- the `common_name` field can be `Grape`;

- the `organism_biocview` field has to be `Vitis_vinifera` (see this link).

Forge the target package from the seed file:

```
echo "date; echo \"library(BSgenome); forgeBSgenomeDataPkg(\\\"IGGP12Xv2_seed.txt\\\
    "); sessionInfo()\" | R --vanilla; date" | qsub -cwd -j y -V -N forge_BSgenome -
    q normal.q
```

Build the package and check it:

```
echo "date; R CMD build BSgenome.Vvinifera.URGI.IGGP12Xv2; date" | qsub -cwd -j y -V
     -N build_BSgenome -q normal.q
echo "date; R CMD check BSgenome.Vvinifera.URGI.IGGP12Xv2_0.1.tar.gz; date" | qsub -
    cwd -j y -V -N check_BSgenome -q normal.q
```

The target package is now ready to be installed:

```
R CMD INSTALL BSgenome.Vvinifera.URGI.IGGP12Xv2_0.1.tar.gz
```

A.-F. Adam-Blondon (part of IGGP) and other colleagues from INRA gave positive feedback. I hence sent the package to the Bioconductor team. I hope it will soon be available for download via `biocLite()`.


### 3.4.2   BSgenome IGGP12Xv0 package

`http://bioconductor.org/packages/release/bioc/html/BSgenome.html`

Retrieve the sequence data from URGI:

```
cd results/
mkdir -p make_BSgenome_IGGP12Xv0
cd make_BSgenome_IGGP12Xv0/
```

```
ln -s ../../data/urgi/VV_chr12x.fsa.gz .
```

TODO: Prepare the seed file (`IGGP12Xv0_seed.txt`) by editing the one used for `IGGP12Xv2`.

### 3.4.3   TxDb IGGP12Xv0 package from NCBI annotations

http://www.bioconductor.org/packages/release/bioc/html/GenomicFeatures.html

TODO: use `makeTxDbFromGFF()`